

Name: Sumanasekara W.K.G.G.
Index: 190610E

```
In [ ]: import numpy as np
import tensorflow as tf
from tensorflow import keras
from keras import datasets, layers, models
```

```
In [ ]: (train_images, train_labels), (test_images, test_labels) = datasets.mnist.load_data()
class_names = [str(i) for i in range(10)]

padding = tf.constant(((0, 0), (2, 2), (2, 2)))
train_images = tf.pad(train_images, padding, constant_values=0)
test_images = tf.pad(test_images, padding, constant_values=0)

print("train_images.shape:", train_images.shape)
print("train_labels.shape:", train_labels.shape)
print("test_images.shape:", test_images.shape)
print("test_labels.shape:", test_labels.shape)

train_images = tf.dtypes.cast(train_images, tf.float32)
test_images = tf.dtypes.cast(test_images, tf.float32)
train_images, test_images = train_images[...]/255.0, test_images[...]/255.0

train_images.shape: (60000, 32, 32)
train_labels.shape: (60000,)
test_images.shape: (10000, 32, 32)
test_labels.shape: (10000,)
```

```
In [ ]: model = models.Sequential()
model.add(layers.Convolution2D(6, (5, 5), activation='relu', input_shape=(32, 32, 1)))
model.add(layers.AveragePooling2D((2, 2)))
model.add(layers.Conv2D(16, (5, 5), activation='relu'))
model.add(layers.AveragePooling2D((2, 2)))
model.add(layers.Flatten())
model.add(layers.Dense(120, activation='relu'))
model.add(layers.Dense(84, activation='relu'))
model.add(layers.Dense(10))

model.compile(optimizer=keras.optimizers.Adam(), loss=keras.losses.SparseCategoricalCrossentropy)
print(model.summary())
```

Model: "sequential_15"

Layer (type)	Output Shape	Param #
conv2d_26 (Conv2D)	(None, 28, 28, 6)	156
average_pooling2d_18 (AveragePooling2D)	(None, 14, 14, 6)	0
conv2d_27 (Conv2D)	(None, 10, 10, 16)	2416
average_pooling2d_19 (AveragePooling2D)	(None, 5, 5, 16)	0
flatten_7 (Flatten)	(None, 400)	0
dense_16 (Dense)	(None, 120)	48120
dense_17 (Dense)	(None, 84)	10164
dense_18 (Dense)	(None, 10)	850
Total params: 61,706		
Trainable params: 61,706		
Non-trainable params: 0		
None		

```
In [ ]: (train_images, train_labels), (test_images, test_labels) = datasets.cifar10.load_data(
class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse']

train_images, test_images = train_images / 255.0, test_images / 255.0

print("train_images.shape:", train_images.shape)
print("train_labels.shape:", train_labels.shape)
print("test_images.shape:", test_images.shape)
print("test_labels.shape:", test_labels.shape)

train_images.shape: (50000, 32, 32, 3)
train_labels.shape: (50000, 1)
test_images.shape: (10000, 32, 32, 3)
test_labels.shape: (10000, 1)
```

```
In [ ]: model = models.Sequential()
model.add(layers.Conv2D(32, (5, 5), activation='relu', input_shape=(32, 32, 3)))
model.add(layers.MaxPool2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPool2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.Flatten())
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(10))

model.compile(optimizer=keras.optimizers.Adam(learning_rate=0.001), loss=keras.losses.categorical_crossentropy)
print(model.summary())
```

Model: "sequential_17"

Layer (type)	Output Shape	Param #
conv2d_31 (Conv2D)	(None, 28, 28, 32)	2432
max_pooling2d_6 (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_32 (Conv2D)	(None, 12, 12, 64)	18496
max_pooling2d_7 (MaxPooling2D)	(None, 6, 6, 64)	0
conv2d_33 (Conv2D)	(None, 4, 4, 128)	73856
flatten_9 (Flatten)	(None, 2048)	0
dense_21 (Dense)	(None, 64)	131136
dense_22 (Dense)	(None, 10)	650
Total params: 226,570		
Trainable params: 226,570		
Non-trainable params: 0		
None		

```
In [ ]: model.fit(train_images, train_labels, epochs=5)
test_loss, test_acc = model.evaluate(test_images, test_labels, verbose=2)
```

Epoch 1/5
1563/1563 [=====] - 48s 31ms/step - loss: 1.5140 - accuracy: 0.4488
Epoch 2/5
1563/1563 [=====] - 51s 33ms/step - loss: 1.1575 - accuracy: 0.5914
Epoch 3/5
1563/1563 [=====] - 44s 28ms/step - loss: 0.9839 - accuracy: 0.6542
Epoch 4/5
1563/1563 [=====] - 44s 28ms/step - loss: 0.8724 - accuracy: 0.6929
Epoch 5/5
1563/1563 [=====] - 46s 30ms/step - loss: 0.7826 - accuracy: 0.7265
313/313 - 2s - loss: 0.8982 - accuracy: 0.6881 - 2s/epoch - 8ms/step

```
In [ ]: (train_images, train_labels), (test_images, test_labels) = datasets.mnist.load_data()
class_names = [str(i) for i in range(10)]

paddings = tf.constant(((0, 0), (2, 2), (2, 2)))
train_images = tf.pad(train_images, paddings, constant_values=0)
test_images = tf.pad(test_images, paddings, constant_values=0)

print("train_images.shape:", train_images.shape)
print("train_labels.shape:", train_labels.shape)
print("test_images.shape:", test_images.shape)
print("test_labels.shape:", test_labels.shape)
```

```
train_images = tf.dtypes.cast(train_images, tf.float32)
test_images = tf.dtypes.cast(test_images, tf.float32)
train_images, test_images = train_images[..., np.newaxis]/255.0, test_images[..., np.newaxis]/255.0

model = models.Sequential()
model.add(layers.Convolution2D(32, (3, 3), activation='relu', input_shape=(32, 32, 1)))
model.add(layers.MaxPool2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPool2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.Flatten())
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(10))

model.compile(optimizer=keras.optimizers.Adam(), loss=keras.losses.SparseCategoricalCrossentropy, metrics=['accuracy'])
print(model.summary())

model.fit(train_images, train_labels, epochs=2)
test_loss, test_acc = model.evaluate(test_images, test_labels, verbose=2)
model.save_weights('saved_weights/')
```

```

train_images.shape: (60000, 32, 32)
train_labels.shape: (60000,)
test_images.shape: (10000, 32, 32)
test_labels.shape: (10000,)
Model: "sequential_18"

```

Layer (type)	Output Shape	Param #
conv2d_34 (Conv2D)	(None, 30, 30, 32)	320
max_pooling2d_8 (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_35 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_9 (MaxPooling2D)	(None, 6, 6, 64)	0
conv2d_36 (Conv2D)	(None, 4, 4, 64)	36928
flatten_10 (Flatten)	(None, 1024)	0
dense_23 (Dense)	(None, 64)	65600
dense_24 (Dense)	(None, 10)	650

=====
 Total params: 121,994
 Trainable params: 121,994
 Non-trainable params: 0

None

Epoch 1/2

1875/1875 [=====] - 46s 24ms/step - loss: 0.1306 - accuracy: 0.9595

Epoch 2/2

1875/1875 [=====] - 47s 25ms/step - loss: 0.0407 - accuracy: 0.9869

313/313 - 2s - loss: 0.0295 - accuracy: 0.9905 - 2s/epoch - 7ms/step

```

In [ ]: model_lw = models.Sequential()
model_lw.add(layers.Convolution2D(32, (3, 3), activation='relu', input_shape=(32, 32, 1)))
model_lw.add(layers.MaxPool2D((2, 2)))
model_lw.add(layers.Conv2D(64, (3, 3), activation='relu'))
model_lw.add(layers.MaxPool2D((2, 2)))
model_lw.add(layers.Conv2D(64, (3, 3), activation='relu'))
model_lw.add(layers.Flatten())
model_lw.add(layers.Dense(64, activation='relu'))
model_lw.add(layers.Dense(10))

model_lw.compile(optimizer=keras.optimizers.Adam(), loss=keras.losses.SparseCategoricalCrossentropy)
print(model_lw.summary())

model_lw.load_weights('saved_weights/')

model_lw.fit(train_images, train_labels, epochs=2)
test_loss, test_acc = model_lw.evaluate(test_images, test_labels, verbose=2)

model_lw.save('saved_model/')

```

Model: "sequential_21"

Layer (type)	Output Shape	Param #
conv2d_43 (Conv2D)	(None, 30, 30, 32)	320
max_pooling2d_14 (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_44 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_15 (MaxPooling2D)	(None, 6, 6, 64)	0
conv2d_45 (Conv2D)	(None, 4, 4, 64)	36928
flatten_13 (Flatten)	(None, 1024)	0
dense_29 (Dense)	(None, 64)	65600
dense_30 (Dense)	(None, 10)	650

Total params: 121,994

Trainable params: 121,994

Non-trainable params: 0

None

Epoch 1/2

1875/1875 [=====] - 46s 24ms/step - loss: 0.0303 - accuracy: 0.9904

Epoch 2/2

1875/1875 [=====] - 44s 23ms/step - loss: 0.0228 - accuracy: 0.9933

313/313 - 3s - loss: 0.0286 - accuracy: 0.9917 - 3s/epoch - 8ms/step

INFO:tensorflow:Assets written to: saved_model/assets

```
In [ ]: model_ld = models.load_model('saved_model/')
        print(model_ld.summary())
```

Model: "sequential_21"

Layer (type)	Output Shape	Param #
conv2d_43 (Conv2D)	(None, 30, 30, 32)	320
max_pooling2d_14 (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_44 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_15 (MaxPooling2D)	(None, 6, 6, 64)	0
conv2d_45 (Conv2D)	(None, 4, 4, 64)	36928
flatten_13 (Flatten)	(None, 1024)	0
dense_29 (Dense)	(None, 64)	65600
dense_30 (Dense)	(None, 10)	650
Total params: 121,994		
Trainable params: 121,994		
Non-trainable params: 0		
None		

```
In [ ]: base_inputs = model_ld.layers[0].input
base_outputs = model_ld.layers[-2].output
output = layers.Dense(10)(base_outputs)

new_model = keras.Model(inputs=base_inputs, outputs=output)
new_model.compile(optimizer=keras.optimizers.Adam(), loss=tf.losses.SparseCategoricalCrossentropy)
print(new_model.summary())

new_model.fit(train_images, train_labels, epochs=3, verbose=2)
test_loss, test_acc = model_lw.evaluate(test_images, test_labels, verbose=2)
```

Model: "model_4"

Layer (type)	Output Shape	Param #
conv2d_43_input (InputLayer)	[(None, 32, 32, 1)]	0
conv2d_43 (Conv2D)	(None, 30, 30, 32)	320
max_pooling2d_14 (MaxPoolin g2D)	(None, 15, 15, 32)	0
conv2d_44 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_15 (MaxPoolin g2D)	(None, 6, 6, 64)	0
conv2d_45 (Conv2D)	(None, 4, 4, 64)	36928
flatten_13 (Flatten)	(None, 1024)	0
dense_29 (Dense)	(None, 64)	65600
dense_39 (Dense)	(None, 10)	650

=====
Total params: 121,994
Trainable params: 121,994
Non-trainable params: 0

None
Epoch 1/3
1875/1875 - 43s - loss: 0.0997 - accuracy: 0.9747 - 43s/epoch - 23ms/step
Epoch 2/3
1875/1875 - 44s - loss: 0.0140 - accuracy: 0.9955 - 44s/epoch - 24ms/step
Epoch 3/3
1875/1875 - 42s - loss: 0.0102 - accuracy: 0.9967 - 42s/epoch - 22ms/step
313/313 - 2s - loss: 0.0286 - accuracy: 0.9917 - 2s/epoch - 7ms/step

```
In [ ]: model_for_tl = models.load_model('saved_model/')
model_for_tl.trainable = False
for layer in model_for_tl.layers:
    assert layer.trainable == False

base_inputs = model_for_tl.layers[0].input
base_outputs = model_for_tl.layers[-2].output
output = layers.Dense(10)(base_outputs)

new_model = keras.Model(inputs=base_inputs, outputs=output)
new_model.compile(optimizer=keras.optimizers.Adam(), loss=tf.losses.SparseCategoricalCrossentropy)
print(new_model.summary())

new_model.fit(train_images, train_labels, epochs=3, verbose=2)
test_loss, test_acc = model_lw.evaluate(test_images, test_labels, verbose=2)
```


Model: "model_5"

Layer (type)	Output Shape	Param #
conv2d_43_input (InputLayer)	[(None, 32, 32, 1)]	0
conv2d_43 (Conv2D)	(None, 30, 30, 32)	320
max_pooling2d_14 (MaxPoolin g2D)	(None, 15, 15, 32)	0
conv2d_44 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_15 (MaxPoolin g2D)	(None, 6, 6, 64)	0
conv2d_45 (Conv2D)	(None, 4, 4, 64)	36928
flatten_13 (Flatten)	(None, 1024)	0
dense_29 (Dense)	(None, 64)	65600
dense_40 (Dense)	(None, 10)	650

```

=====
Total params: 121,994
Trainable params: 650
Non-trainable params: 121,344

```

None

Epoch 1/3

1875/1875 - 12s - loss: 0.1938 - accuracy: 0.9566 - 12s/epoch - 6ms/step

Epoch 2/3

1875/1875 - 12s - loss: 0.0146 - accuracy: 0.9961 - 12s/epoch - 6ms/step

Epoch 3/3

1875/1875 - 12s - loss: 0.0106 - accuracy: 0.9969 - 12s/epoch - 6ms/step

313/313 - 2s - loss: 0.0286 - accuracy: 0.9917 - 2s/epoch - 7ms/step

```

In [ ]: model_rs = models.Sequential()
        model_rs.add(keras.applications.ResNet50V2(include_top=True, weights='imagenet', input
                                                    input_shape=None, pooling=None, classes=1000,
                                                    classifier_activation='softmax'))

        model_rs.trainable = False
        for layer in model_rs.layers:
            assert layer.trainable == False

```